

PyTorch-to-Quantized-C Compiler

Compiling Neural Networks for Microcontrollers

Ashitabh Misra, Joydeep Bhattacharyya, Tarek Abdelzaher

January 9, 2026

Outline

- 1 Introduction & Motivation
- 2 Tutorial & Use Cases
- 3 Inner Workings
- 4 Writing Optimization Passes
- 5 Testing & Results
- 6 Future Directions

The Problem

Deploying quantized neural networks on microcontrollers is challenging.

Existing tools are inflexible and difficult to extend:

- **TFLite Micro:** Large binary size for microcontrollers usecases, not designed to be extended. As a result, extremely difficult to add new quantization techniques.

Deploying quantized neural networks on microcontrollers is challenging.

Existing tools are inflexible and difficult to extend:

- **TFLite Micro:** Large binary size for microcontrollers usecases, not designed to be extended. As a result, extremely difficult to add new quantization techniques.
- **Glow:** Complex setup, designed for research not real-world embedded use, hard to add new features.

Comparison with Existing Tools

Feature	Tiny-NN-in-C	TFLite Micro	Glow

Comparison with Existing Tools

Feature	Tiny-NN-in-C	TFLite Micro	Glow
Framework	PyTorch	TensorFlow	ONNX/PyTorch*

Comparison with Existing Tools

Feature	Tiny-NN-in-C	TFLite Micro	Glow
Framework Output	PyTorch Standalone C	TensorFlow .tflite + Interp.	ONNX/PyTorch* Compiled binary

Comparison with Existing Tools

Feature	Tiny-NN-in-C	TFLite Micro	Glow
Framework	PyTorch	TensorFlow	ONNX/PyTorch*
Output	Standalone C	.tflite + Interp.	Compiled binary
Code Readability			

Comparison with Existing Tools

Feature	Tiny-NN-in-C	TFLite Micro	Glow
Framework	PyTorch	TensorFlow	ONNX/PyTorch*
Output	Standalone C	.tflite + Interp.	Compiled binary
Code Readability	✓	✗ Binary	✗ Opt. IR
Setup Complexity			

Comparison with Existing Tools

Feature	Tiny-NN-in-C	TFLite Micro	Glow
Framework	PyTorch	TensorFlow	ONNX/PyTorch*
Output	Standalone C	.tflite + Interp.	Compiled binary
Code Readability	✓	✗ Binary	✗ Opt. IR
Setup Complexity	Simple (pip)	Medium	Complex
Binary Size			

Comparison with Existing Tools

Feature	Tiny-NN-in-C	TFLite Micro	Glow
Framework	PyTorch	TensorFlow	ONNX/PyTorch*
Output	Standalone C	.tflite + Interp.	Compiled binary
Code Readability	✓	✗ Binary	✗ Opt. IR
Setup Complexity	Simple (pip)	Medium	Complex
Binary Size	~150KB	~390KB+	Varies
Quantization			

Comparison with Existing Tools

Feature	Tiny-NN-in-C	TFLite Micro	Glow
Framework	PyTorch	TensorFlow	ONNX/PyTorch*
Output	Standalone C	.tflite + Interp.	Compiled binary
Code Readability	✓	✗ Binary	✗ Opt. IR
Setup Complexity	Simple (pip)	Medium	Complex
Binary Size	~150KB	~390KB+	Varies
Quantization	int8/16, mixed	int8	int8
Custom Rules			

Comparison with Existing Tools

Feature	Tiny-NN-in-C	TFLite Micro	Glow
Framework	PyTorch	TensorFlow	ONNX/PyTorch*
Output	Standalone C	.tflite + Interp.	Compiled binary
Code Readability	✓	✗ Binary	✗ Opt. IR
Setup Complexity	Simple (pip)	Medium	Complex
Binary Size	~150KB	~390KB+	Varies
Quantization	int8/16, mixed	int8	int8
Custom Rules			
Extensibility			

Comparison with Existing Tools

Feature	Tiny-NN-in-C	TFLite Micro	Glow
Framework	PyTorch	TensorFlow	ONNX/PyTorch*
Output	Standalone C	.tflite + Interp.	Compiled binary
Code Readability	✓	✗ Binary	✗ Opt. IR
Setup Complexity	Simple (pip)	Medium	Complex
Binary Size	~150KB	~390KB+	Varies
Quantization	int8/16, mixed	int8	int8
Custom Rules	✓ Regex	✗ Per-layer	✗ Manual
Extensibility			
Transparency			

Comparison with Existing Tools

Feature	Tiny-NN-in-C	TFLite Micro	Glow
Framework	PyTorch	TensorFlow	ONNX/PyTorch*
Output	Standalone C	.tflite + Interp.	Compiled binary
Code Readability	✓	✗ Binary	✗ Opt. IR
Setup Complexity	Simple (pip)	Medium	Complex
Binary Size	~150KB	~390KB+	Varies
Quantization	int8/16, mixed	int8	int8
Custom Rules	✓ Regex	✗ Per-layer	✗ Manual
Extensibility			
Transparency			

Comparison with Existing Tools

Feature	Tiny-NN-in-C	TFLite Micro	Glow
Framework	PyTorch	TensorFlow	ONNX/PyTorch*
Output	Standalone C	.tflite + Interp.	Compiled binary
Code Readability	✓	✗ Binary	✗ Opt. IR
Setup Complexity	Simple (pip)	Medium	Complex
Binary Size	~150KB	~390KB+	Varies
Quantization	int8/16, mixed	int8	int8
Custom Rules	✓ Regex	✗ Per-layer	✗ Manual
Extensibility	✓ Easy	✗ C++	✗ LLVM
Transparency			

*Glow development has slowed; PyTorch support via ONNX conversion

Comparison with Existing Tools

Feature	Tiny-NN-in-C	TFLite Micro	Glow
Framework	PyTorch	TensorFlow	ONNX/PyTorch*
Output	Standalone C	.tflite + Interp.	Compiled binary
Code Readability	✓	✗ Binary	✗ Opt. IR
Setup Complexity	Simple (pip)	Medium	Complex
Binary Size	~150KB	~390KB+	Varies
Quantization	int8/16, mixed	int8	int8
Custom Rules	✓ Regex	✗ Per-layer	✗ Manual
Extensibility	✓ Easy	✗ C++	✗ LLVM
Transparency			

*Glow development has slowed; PyTorch support via ONNX conversion

Comparison with Existing Tools

Feature	Tiny-NN-in-C	TFLite Micro	Glow
Framework	PyTorch	TensorFlow	ONNX/PyTorch*
Output	Standalone C	.tflite + Interp.	Compiled binary
Code Readability	✓	✗ Binary	✗ Opt. IR
Setup Complexity	Simple (pip)	Medium	Complex
Binary Size	~150KB	~390KB+	Varies
Quantization	int8/16, mixed	int8	int8
Custom Rules	✓ Regex	✗ Per-layer	✗ Manual
Extensibility	✓ Easy	✗ C++	✗ LLVM
Transparency	✓ See code	✗ Black box	✗ Black box

*Glow development has slowed; PyTorch support via ONNX conversion


Architecture Comparison

TensorFlow Micro:

TensorFlow
Model


Architecture Comparison

TensorFlow Micro:




Architecture Comparison

TensorFlow Micro:



Architecture Comparison

TFLite Micro:




Architecture Comparison

TFLite Micro:



Architecture Comparison

TFLite Micro:




Architecture Comparison


TFLite Micro:




Glow:



Our Approach: Tiny-NN-in-C




Our Approach: Tiny-NN-in-C




Our Approach: Tiny-NN-in-C




Our Approach: Tiny-NN-in-C




Our Approach: Tiny-NN-in-C



Our Approach: Tiny-NN-in-C



Our Approach: Tiny-NN-in-C



Key Insight

We achieve **transparency** and **simplicity** with minimal performance penalty.

Quick Start: The Core API

Three simple steps to compile a quantized model:

Step 1: Compile to C (float)

```
from src.pytorch_to_c.compiler import compile_model  
  
ir_graph = compile_model(model, example_input, output_dir="out/")
```

Quick Start: The Core API

Three simple steps to compile a quantized model:

Step 1: Compile to C (float)

```
from src.pytorch_to_c.compiler import compile_model  
  
ir_graph = compile_model(model, example_input, output_dir="out/")
```

Step 2: Define quantization rules

```
from src.pytorch_to_c.quantization import StaticQuantRule  
  
rules = [StaticQuantRule(pattern=r'fc.*', dtype='int8', ...)]
```

Quick Start: The Core API

Three simple steps to compile a quantized model:

Step 1: Compile to C (float)

```
from src.pytorch_to_c.compiler import compile_model  
  
ir_graph = compile_model(model, example_input, output_dir="out/")
```

Step 2: Define quantization rules

```
from src.pytorch_to_c.quantization import StaticQuantRule  
  
rules = [StaticQuantRule(pattern=r'fc.*', dtype='int8', ...)]
```

Step 3: Apply quantization and generate

```
transform = QuantizationTransform(rules)  
quant_ir = transform.apply(ir_graph)  
CPrinter(quant_ir).generate_all("out/")
```

What is This?

A PyTorch-to-C compiler for microcontrollers

- Takes PyTorch models → Generates standalone C code

What is This?

A PyTorch-to-C compiler for microcontrollers

- Takes PyTorch models → Generates standalone C code
- Supports quantization (int8, int16) for embedded deployment

A PyTorch-to-C compiler for microcontrollers

- Takes PyTorch models → Generates standalone C code
- Supports quantization (int8, int16) for embedded deployment
- Extensible rule-based system for selective quantization

A PyTorch-to-C compiler for microcontrollers

- Takes PyTorch models → Generates standalone C code
- Supports quantization (int8, int16) for embedded deployment
- Extensible rule-based system for selective quantization
- IR-based architecture enables optimization passes

What is This?

A PyTorch-to-C compiler for microcontrollers

- Takes PyTorch models → Generates standalone C code
- Supports quantization (int8, int16) for embedded deployment
- Extensible rule-based system for selective quantization
- IR-based architecture enables optimization passes

```
python examples/tiny_resnet.py
```

Simple MLP Example

Model: 2-layer MLP (Linear → ReLU → Linear)

```
class SimpleMLP(nn.Module):
    def __init__(self, input_size=784, hidden_size=128, output_size=10):
        super().__init__()
```

Simple MLP Example

Model: 2-layer MLP (Linear → ReLU → Linear)

```
class SimpleMLP(nn.Module):
    def __init__(self, input_size=784, hidden_size=128, output_size=10):
        super().__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(hidden_size, output_size)
```

Simple MLP Example

Model: 2-layer MLP (Linear → ReLU → Linear)

```
class SimpleMLP(nn.Module):
    def __init__(self, input_size=784, hidden_size=128, output_size=10):
        super().__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(hidden_size, output_size)

    def forward(self, x):
        x = self.fc1(x)
        x = self.relu(x)
        x = self.fc2(x)
        return x
```

Simple MLP Example

Compile to C:

```
model = SimpleMLP(input_size=784, hidden_size=128, output_size=10)
example_input = torch.randn(1, 784)

ir_graph = compile_model(
    model=model,
    example_input=example_input,
    output_dir="generated"
)
```

Simple MLP Example: Generated IR Graph

IR Graph structure:

```
IRGraph:  
  Inputs: ['x']  
  Outputs: ['fc2']  
  Parameters: ['fc1_weight', 'fc1_bias', 'fc2_weight', 'fc2_bias']  
  Nodes:  
    x [input]  
      inputs: []  
      users: [fc1]  
      shape: (1, 784), dtype: float32  
    fc1 [linear]  
      inputs: [x]  
      users: [relu]  
      shape: (1, 128), dtype: float32  
    relu [relu]  
      inputs: [fc1]  
      users: [fc2]  
      shape: (1, 128), dtype: float32  
    fc2 [linear]  
      inputs: [relu]  
      users: []  
      shape: (1, 10), dtype: float32
```

ResNet-Style Model (Skip Connections)

Architecture: Conv → BatchNorm → ReLU → Skip → Pool → FC

```
class ResNetBlock(nn.Module):
    def __init__(self, channels):
        super().__init__()
        self.conv1 = nn.Conv2d(channels, channels, 3, padding=1)
        self.bn1 = nn.BatchNorm2d(channels)
```

ResNet-Style Model (Skip Connections)

Architecture: Conv → BatchNorm → ReLU → Skip → Pool → FC

```
class ResNetBlock(nn.Module):
    def __init__(self, channels):
        super().__init__()
        self.conv1 = nn.Conv2d(channels, channels, 3, padding=1)
        self.bn1 = nn.BatchNorm2d(channels)

    def forward(self, x):
        identity = x
        out = torch.relu(self.bn1(self.conv1(x)))
        return out + identity
```

ResNet-Style Model (Skip Connections)

TinyResNet: Using ResNetBlock with skip connections

```
class TinyResNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv_init = nn.Conv2d(3, 32, 3, padding=1)
        self.block1 = ResNetBlock(32)
        self.fc = nn.Linear(32, 10)

    def forward(self, x):
        x = torch.relu(self.conv_init(x))
        x = self.block1(x)
        x = x.mean(dim=[2, 3])
        return self.fc(x)
```

Static Quantization

Key idea: Scales are known at compile time (from calibration)

Step 1: Define quantization rules

```
from src.pytorch_to_c.quantization import StaticQuantRule

rules = [
    StaticQuantRule(
        pattern=r'fc.*',
        dtype='int8',
```

Static Quantization

Key idea: Scales are known at compile time (from calibration)

Step 1: Define quantization rules

```
from src.pytorch_to_c.quantization import StaticQuantRule

rules = [
    StaticQuantRule(
        pattern=r'fc.*',
        dtype='int8',
        input_scale=0.01, weight_scale=0.01, output_scale=0.01
    )
]
```

Static Quantization

Key idea: Scales are known at compile time (from calibration)

Step 2: Apply to IR graph

```
ir_graph = compile_model(model, example_input, return_ir=True)
transform = QuantizationTransform(rules)
quant_ir = transform.apply(ir_graph)
```

Dynamic Quantization

Key idea: Input scale computed at runtime from actual data

Define dynamic rule (no calibration needed!):

```
from src.pytorch_to_c.quantization import DynamicQuantRuleMinMaxPerTensor

rules = [
    DynamicQuantRuleMinMaxPerTensor(
        pattern=r'fc.*',
        dtype='int8'
    )
]
```

Dynamic Quantization

Key idea: Input scale computed at runtime from actual data

Generated C code computes scale dynamically:

```
float scale = compute_dynamic_scale_int8(input, 784);
quantize_float_to_int8(input, 784, scale, 0, buf_quantized);
```

Advantage

No calibration dataset required—scales derived from weight statistics.

Mixed Precision Quantization

Selective quantization: Different rules for different layers

```
rules = [  
    StaticQuantRule(pattern=r'.*encoder.*', dtype='int8', ...),  
]
```

Mixed Precision Quantization

Selective quantization: Different rules for different layers

```
rules = [
    StaticQuantRule(pattern=r'.*encoder.*', dtype='int8', ...),
    StaticQuantRule(pattern=r'.*output.*', dtype='int16', ...),
]
```

Correctness Results (On TinyResNet model)

Comparing our implementation with Pytorch Float32 model.

Configuration	Max Error	Memory Savings
---------------	-----------	----------------

Correctness Results (On TinyResNet model)

Comparing our implementation with Pytorch Float32 model.

Configuration	Max Error	Memory Savings
Float32 (baseline)	1.19×10^{-7}	—

Correctness Results (On TinyResNet model)

Comparing our implementation with Pytorch Float32 model.

Configuration	Max Error	Memory Savings
Float32 (baseline)	1.19×10^{-7}	—
Static int16	0.07%	2×

Correctness Results (On TinyResNet model)

Comparing our implementation with Pytorch Float32 model.

Configuration	Max Error	Memory Savings
Float32 (baseline)	1.19×10^{-7}	—
Static int16	0.07%	2×
Static int8	1.63%	4×

Correctness Results (On TinyResNet model)

Comparing our implementation with Pytorch Float32 model.

Configuration	Max Error	Memory Savings
Float32 (baseline)	1.19×10^{-7}	—
Static int16	0.07%	2×
Static int8	1.63%	4×
Dynamic int8	2.95%	4×

Correctness Results (On TinyResNet model)

Comparing our implementation with Pytorch Float32 model.

Configuration	Max Error	Memory Savings
Float32 (baseline)	1.19×10^{-7}	—
Static int16	0.07%	2×
Static int8	1.63%	4×
Dynamic int8	2.95%	4×

Takeaway

Choose precision based on your accuracy requirements and memory constraints.

Optimization Pass: FuseDequantQuant

Problem: Consecutive quantized layers have redundant conversions

```
BEFORE: fc1(int8) -> dequant(float) -> quant(int8) -> fc2(int8)
```

Optimization Pass: FuseDequantQuant

Problem: Consecutive quantized layers have redundant conversions

```
BEFORE: fc1(int8) -> dequant(float) -> quant(int8) -> fc2(int8)  
AFTER:  fc1(int8) -> fc2(int8)
```

Optimization Pass: FuseDequantQuant

Apply the optimization pass:

```
from src.passes import FuseDequantQuantPass  
  
fuse_pass = FuseDequantQuantPass()  
optimized_ir = fuse_pass.apply(quant_ir)
```

Optimization Pass: FuseDequantQuant


Apply the optimization pass:

```
from src.passes import FuseDequantQuantPass  
  
fuse_pass = FuseDequantQuantPass()  
optimized_ir = fuse_pass.apply(quant_ir)
```

Result

When scales match, fusion is **bit-identical**: Max error = 0.00e+00

Architecture Overview



Key files:

- `frontend/fx_tracer.py` – torch.fx tracing
- `lowering/lower.py` – FX → IR conversion
- `ir/graph.py` – IR graph structure
- `codegen/c_printer.py` – C code generation

IR Node Structure

Base class: src/pytorch_to_c/ir/node.py

```
class IRNode:
    def __init__(self, name, op_type, inputs=None, dtype='float32'):
        self.name = name          # Unique identifier
        self.op_type = op_type      # 'linear', 'conv2d', 'relu', etc.
        self.inputs = inputs or [] # List of input IRNodes
        self.users = []            # Nodes that use this output
        self.dtype = dtype          # 'float32', 'int8', 'int16'
        self.output_shape = None    # Inferred shape
        self.metadata = {}          # Extra info (weights, etc.)

    def get_c_dtype(self) -> str:
        """Returns C type: 'float', 'int8_t', 'int16_t'"""

```

Graph is doubly-linked: inputs + users enable traversal in both directions

IR Graph Structure

File: src/pytorch_to_c/ir/graph.py

```
class IRGraph:
    def __init__(self):
        self.nodes = []
        self.inputs = []
        self.outputs = []
        self.parameters = {}

    def print_graph(self):
        """Pretty-print graph"""

```

Example IR output:

```
x [input] dtype=float32
    inputs: []
    users: [fc1]
fc1 [linear] dtype=float32
    inputs: [x]
    users: [relu]
    shape: (1, 8)
```

Quantization Rule System

Base class: `src/pytorch_to_c/quantization/rules.py`

```
class QuantRule(ABC):
    def __init__(self, pattern: str, dtype: str):
        self.pattern = re.compile(pattern) # Regex for layer names
        self.dtype = dtype # 'int8' or 'int16'

    def matches(self, node: IRNode) -> bool:
        """Check if rule applies to this node"""
        return self.pattern.search(node.name) is not None

    @abstractmethod
    def create_quant_node(self, float_node: IRNode) -> QuantIRNode:
        """Create quantized version of the node"""


```

Rule matching: First matching rule wins

```
RuleMatcher([rule1, rule2, rule3]) # Checked in order
```

Writing a Custom Rule

Example: Static Quantization Rule

```
class StaticQuantRule(QuantRule):
    def __init__(self, pattern, dtype,
                 input_scale, weight_scale, output_scale, ...):
        super().__init__(pattern, dtype)
        self.input_scale = input_scale
        # ... store all scales/offsets
```

Writing a Custom Rule

Example: Static Quantization Rule

```
class StaticQuantRule(QuantRule):
    def __init__(self, pattern, dtype,
                 input_scale, weight_scale, output_scale, ...):
        super().__init__(pattern, dtype)
        self.input_scale = input_scale

    def create_quant_node(self, float_node: IRNode) -> QuantIRNode:
        if float_node.op_type == 'linear':
            return StaticQuantLinearNode(
                name=float_node.name, dtype=self.dtype,
                input_scale=self.input_scale, ...
            )
        elif float_node.op_type == 'conv2d':
            return StaticQuantConv2dNode(...)
```

Quantized Node Structure

Base class: `src/pytorch_to_c/ir/quant_node.py`

```
class QuantIRNode(IRNode):
    """Base class for all quantized operations"""

    @abstractmethod
    def get_pre_nodes(self) -> List[IRNode]:
        """Nodes to insert BEFORE this op (e.g., QuantizeNode)"""

    @abstractmethod
    def get_post_nodes(self) -> List[IRNode]:
        """Nodes to insert AFTER this op (e.g., DequantizeNode)"""

    @abstractmethod
    def generate_c_code(self, c_printer) -> str:
        """Generate the C code for this operation"""
```

Key insight: Each QuantIRNode controls its own conversion nodes!

StaticQuantLinearNode Implementation

File: src/pytorch_to_c/quantization/ops/quant_linear.py

```
class StaticQuantLinearNode(QuantIRNode):
    def __init__(self, name, dtype, input_scale, weight_scale, ...):
        super().__init__(name, 'linear', dtype=dtype)
        self.input_scale = input_scale
        self.weight_scale = weight_scale
        self.output_scale = output_scale
```

StaticQuantLinearNode Implementation

File: src/pytorch_to_c/quantization/ops/quant_linear.py

```
def get_pre_nodes(self):
    return [QuantizeNode(
        name=f"{self.name}_input_q",
        dtype=self.dtype,
        scale=self.input_scale,
        offset=self.input_offset
    )]

def get_post_nodes(self):
    return [DequantizeNode(
        name=f"{self.name}_output_dq",
        scale=self.output_scale,
        offset=self.output_offset
    )]
```

StaticQuantLinearNode Implementation

File: src/pytorch_to_c/quantization/ops/quant_linear.py

```
def generate_c_code(self, c_printer):
    return f"dense_int8({{input}}, {{size}}, {{weight}}, " \
           f"{{bias}}, {{out_size}}, {{self.input_scale}f}, " \
           f"{{self.weight_scale}f}, 0, {{output}});"
```

Quantization Transform Pipeline

File: src/pytorch_to_c/quantization/graph_transform.py

```
class QuantizationTransform:
    def __init__(self, rules: List[QuantRule]):
        self.matcher = RuleMatcher(rules)

    def apply(self, ir_graph: IRGraph) -> IRGraph:
        # Step 1: Find nodes matching rules
        nodes_to_quantize = self._find_nodes_to_quantize(ir_graph)

        # Step 2: Replace float nodes with QuantIRNodes
        self._replace_nodes(ir_graph, nodes_to_quantize)

        # Step 3: Insert pre/post nodes (Quantize/Dequantize)
        self._insert_node_controlled_conversions(ir_graph)

        # Step 4: Validate output is float32
        self._validate_float_output(ir_graph)

        # Step 5: Quantize weights
        self._quantize_weights(ir_graph, nodes_to_quantize)

    return ir_graph
```

Adding a New Quantized Operation

Steps to add QuantizedSoftmax:

- ① Create the node class:

```
class StaticQuantSoftmaxNode(QuantIRNode):  
    def get_pre_nodes(self): ...  
    def get_post_nodes(self): ...  
    def generate_c_code(self, c_printer): ...
```

Adding a New Quantized Operation

Steps to add QuantizedSoftmax:

① Create the node class:

```
class StaticQuantSoftmaxNode(QuantIRNode):  
    def get_pre_nodes(self): ...  
    def get_post_nodes(self): ...  
    def generate_c_code(self, c_printer): ...
```

② Add C implementation:

```
void softmax_int8(const int8_t* input, int size,  
                  float scale, int offset, int8_t* output);
```

Adding a New Quantized Operation

Steps to add QuantizedSoftmax:

- ① Create the node class:

```
class StaticQuantSoftmaxNode(QuantIRNode):  
    def get_pre_nodes(self): ...  
    def get_post_nodes(self): ...  
    def generate_c_code(self, c_printer): ...
```

- ② Add C implementation:

```
void softmax_int8(const int8_t* input, int size,  
                  float scale, int offset, int8_t* output);
```

- ③ Update rule to create it
- ④ Export in __init__.py

Pass Infrastructure

Base class: src/passes/base.py

```
class IRPass(ABC):
    def __init__(self, verbose: bool = False):
        self.verbose = verbose
        self.stats = {} # Track optimization statistics

    @abstractmethod
    def apply(self, ir_graph: IRGraph) -> IRGraph:
        """Transform the IR graph"""
        pass

    def get_stats(self) -> Dict[str, Any]:
        """Return statistics about optimizations applied"""
        return self.stats
```

FuseDequantQuantPass Implementation

File: src/passes/fuse_dequant_quant.py

```
class FuseDequantQuantPass(IRPass):
    def apply(self, ir_graph: IRGraph) -> IRGraph:
        # 1. Find fusible pairs
        pairs = self._find_fusible_pairs(ir_graph)

        # 2. Fuse each pair
        for dequant, quant in reversed(pairs):
            self._fuse_pair(ir_graph, dequant, quant)

    return ir_graph
```

FuseDequantQuantPass Implementation

File: src/passes/fuse_dequant_quant.py

```
def _find_fusable_pairs(self, ir_graph):
    pairs = []
    for node in ir_graph.nodes:
        if isinstance(node, DequantizeNode):
            if len(node.users) == 1:
                user = node.users[0]
                if isinstance(user, QuantizeNode):
                    # Check: same dtype AND same scale
                    if (node.inputs[0].dtype == user.dtype and
                        self._scales_match(node, user)):
                        pairs.append((node, user))
    return pairs
```

Graph Rewiring in Passes

Before fusion:

```
A -> dequant -> quant -> B
```

After fusion:

```
A -> B
```

Rewiring code:

```
def _fuse_pair(self, ir_graph, dequant_node, quant_node):
    source = dequant_node.inputs[0]
    targets = quant_node.users.copy()

    # Rewire: source -> targets
    source.users.remove(dequant_node)
    for target in targets:
        target.inputs = [source if x == quant_node else x
                        for x in target.inputs]
        source.users.append(target)

    # Remove nodes from graph
    ir_graph.nodes.remove(dequant_node)
    ir_graph.nodes.remove(quant_node)
```

Writing Your Own Pass

Example: Dead Code Elimination Pass

```
class DeadCodeEliminationPass(IRPass):
    """Remove nodes whose outputs are never used"""

    def apply(self, ir_graph: IRGraph) -> IRGraph:
        changed = True
        while changed:
            changed = False
            for node in ir_graph.nodes.copy():
                # Skip output nodes
                if node in ir_graph.outputs:
                    continue
                # Skip input nodes
                if node in ir_graph.inputs:
                    continue
                # If no users, node is dead
                if len(node.users) == 0:
                    self._remove_node(ir_graph, node)
                    changed = True
        return ir_graph
```

Pass Composition

Chain multiple passes:

```
from src.passes import FuseDequantQuantPass

# Apply quantization
transform = QuantizationTransform(rules)
quant_ir = transform.apply(ir_graph)

# Apply optimization passes
passes = [
    FuseDequantQuantPass(verbose=True),
    # DeadCodeEliminationPass(),
    # ConstantFoldingPass(),
]

optimized_ir = quant_ir
for p in passes:
    optimized_ir = p.apply(optimized_ir)
    print(f"{p.__class__.__name__}: {p.get_stats()}")

# Generate final C code
CPrinter(optimized_ir).generate_all("output/")
```

Test Commands Summary

```
# All tests (64 tests)
python -m pytest test/ -v

# Float model tests
python -m pytest test/test_integration.py -v -s

# Quantization unit tests
python -m pytest test/test_quantization.py -v -s

# Quantization end-to-end tests
python -m pytest test/test_quantization_e2e.py -v -s

# Optimization pass tests
python -m pytest test/test_passes.py -v -s
```

Commands for commit 9362ae0

Key Numerical Checks

Test	Max Error	Test File

Key Numerical Checks

Test	Max Error	Test File
Float MLP	$\sim 10^{-6}$	test_integration.py

Key Numerical Checks

Test	Max Error	Test File
Float MLP	$\sim 10^{-6}$	test_integration.py
Float ResNet	1.19×10^{-7}	test_quantization_e2e.py

Key Numerical Checks

Test	Max Error	Test File
Float MLP	$\sim 10^{-6}$	test_integration.py
Float ResNet	1.19×10^{-7}	test_quantization_e2e.py
Static int8	$\sim 1\text{--}2\%$	test_quantization_e2e.py

Key Numerical Checks

Test	Max Error	Test File
Float MLP	$\sim 10^{-6}$	test_integration.py
Float ResNet	1.19×10^{-7}	test_quantization_e2e.py
Static int8	$\sim 1\text{--}2\%$	test_quantization_e2e.py
Static int16	$\sim 0.1\%$	test_quantization_e2e.py

Key Numerical Checks

Test	Max Error	Test File
Float MLP	$\sim 10^{-6}$	test_integration.py
Float ResNet	1.19×10^{-7}	test_quantization_e2e.py
Static int8	$\sim 1\text{--}2\%$	test_quantization_e2e.py
Static int16	$\sim 0.1\%$	test_quantization_e2e.py
FusePass	0.00e+00	test_passes.py

Key Numerical Checks

Test	Max Error	Test File
Float MLP	$\sim 10^{-6}$	test_integration.py
Float ResNet	1.19×10^{-7}	test_quantization_e2e.py
Static int8	$\sim 1\text{--}2\%$	test_quantization_e2e.py
Static int16	$\sim 0.1\%$	test_quantization_e2e.py
FusePass	0.00e+00	test_passes.py

FusePass is Bit-Identical!

When scales match, the optimization produces exactly the same numerical results.

Project File Structure

```
src/
  pytorch_to_c/
    compiler.py          # Main entry point
    frontend/fx_tracer.py # torch.fx tracing
    lowering/lower.py     # FX -> IR conversion
  ir/
    node.py              # IRNode base class
    graph.py             # IRGraph
    quant_node.py        # QuantIRNode base
    codegen/c_printer.py # C code generation
    quantization/
      rules.py           # QuantRule, StaticQuantRule, etc.
      graph_transform.py # QuantizationTransform
      ops/               # Quantized operation implementations
  passes/
    base.py              # IRPass base class
    fuse_dequant_quant.py # FuseDequantQuantPass
  c_ops/
    nn_ops_float.h       # Float C kernels
    nn_ops_int8.h        # Int8 C kernels
    nn_ops_int16.h        # Int16 C kernels
```

What we built:

- PyTorch → C compiler with torch.fx frontend

What we built:

- PyTorch → C compiler with torch.fx frontend
- Clean IR with doubly-linked node graph

What we built:

- PyTorch → C compiler with torch.fx frontend
- Clean IR with doubly-linked node graph
- Independent from source rule-based quantization

What we built:

- PyTorch → C compiler with torch.fx frontend
- Clean IR with doubly-linked node graph
- Independent from source rule-based quantization
- Support for int8 and int16

What we built:

- PyTorch → C compiler with torch.fx frontend
- Clean IR with doubly-linked node graph
- Independent from source rule-based quantization
- Support for int8 and int16
- Mixed precision quantization

What we built:

- PyTorch → C compiler with torch.fx frontend
- Clean IR with doubly-linked node graph
- Independent from source rule-based quantization
- Support for int8 and int16
- Mixed precision quantization
- Optimization pass infrastructure

What we built:

- PyTorch → C compiler with torch.fx frontend
- Clean IR with doubly-linked node graph
- Independent from source rule-based quantization
- Support for int8 and int16
- Mixed precision quantization
- Optimization pass infrastructure
- FuseDequantQuantPass (bit-identical results!)

Extensibility Points

How to extend the compiler:

Add new quantization:

- New QuantRule subclass
- New QuantIRNode implementation

Add new optimization:

- New IRPass subclass
- Add to pass pipeline

Add new operations:

- Add to lowering rules
- Add C kernel implementation

Add new datatypes:

- Extend get_c_dtype()
- Add C kernels for dtype

Run all 64 tests: `python -m pytest test/ -v`

Extending the compiler for adaptive runtime behavior:

① Resource-Aware Precision Switching

Dynamically switch between quantization precisions at runtime

Extending the compiler for adaptive runtime behavior:

① Resource-Aware Precision Switching

Dynamically switch between quantization precisions at runtime

② Adaptive Operation Implementations

Choose optimal op implementations based on system state

Extending the compiler for adaptive runtime behavior:

① Resource-Aware Precision Switching

Dynamically switch between quantization precisions at runtime

② Adaptive Operation Implementations

Choose optimal op implementations based on system state

③ User-Defined Switching Metrics

Custom policies for runtime adaptation decisions

Resource-Aware Precision Switching

Goal: Enable runtime precision switching based on available resources

Concept:

- Compile model with multiple precision paths (int8, int16, float32)
- Switch between precisions based on battery, thermal, latency constraints

Resource-Aware Precision Switching

Goal: Enable runtime precision switching based on available resources

Concept:

- Compile model with multiple precision paths (int8, int16, float32)
- Switch between precisions based on battery, thermal, latency constraints

Example API:

```
compile_model(  
    model,  
    multi_precision=True,  
    precision_options=['int8', 'int16', 'float32'],  
    switching_policy='battery_aware'  
)
```

Resource-Aware Precision Switching

Goal: Enable runtime precision switching based on available resources

Concept:

- Compile model with multiple precision paths (int8, int16, float32)
- Switch between precisions based on battery, thermal, latency constraints

Example API:

```
compile_model(  
    model,  
    multi_precision=True,  
    precision_options=['int8', 'int16', 'float32'],  
    switching_policy='battery_aware'  
)
```

Generated C code:

```
if (battery_level < 20) {  
    model_forward_int8(input, output); // Fastest, lowest power  
} else if (accuracy_required > 0.95) {  
    model_forward_float32(input, output); // Highest accuracy  
} else {
```

Adaptive Operation Implementations

Goal: Switch between operation implementations based on system state

Use cases:

- Multi-threaded vs. single-threaded based on CPU availability
- SIMD-optimized vs. scalar based on processor features
- Memory-optimized vs. speed-optimized based on RAM pressure

Adaptive Operation Implementations

Goal: Switch between operation implementations based on system state

Use cases:

- Multi-threaded vs. single-threaded based on CPU availability
- SIMD-optimized vs. scalar based on processor features
- Memory-optimized vs. speed-optimized based on RAM pressure

Example: Threading strategy

```
void dense_adaptive(float* input, int size, float* weight,
                     float* output, runtime_context_t* ctx) {
    if (ctx->cpu_cores_available > 2 && size > 1024) {
        dense_multithreaded(input, size, weight, output, ctx);
    } else {
        dense_singlethreaded(input, size, weight, output);
    }
}
```

Adaptive Operation Implementations

Goal: Switch between operation implementations based on system state

Use cases:

- Multi-threaded vs. single-threaded based on CPU availability
- SIMD-optimized vs. scalar based on processor features
- Memory-optimized vs. speed-optimized based on RAM pressure

Example: Threading strategy

```
void dense_adaptive(float* input, int size, float* weight,
                     float* output, runtime_context_t* ctx) {
    if (ctx->cpu_cores_available > 2 && size > 1024) {
        dense_multithreaded(input, size, weight, output, ctx);
    } else {
        dense_singlethreaded(input, size, weight, output);
    }
}
```

Compiler directive:

```
@adaptive_implementation(variants=['single', 'multi', 'simd'])
def forward(...):
```

User-Defined Switching Metrics

Goal: Allow users to define custom policies for runtime adaptation

Custom metric system:

```
class CustomSwitchingPolicy:  
    def should_switch_precision(self, layer_name, current_metrics):  
        latency = current_metrics['latency']  
        accuracy = current_metrics['accuracy']  
  
        if layer_name.startswith('critical_'):  
            return 'float32' if accuracy < 0.99 else 'int16'  
  
        return 'int8' if latency > 50 else 'int16'  
  
    def should_switch_implementation(self, op_name, system_state):  
        if system_state['temperature'] > 70:  
            return 'low_power'  
        return 'high_performance'
```

User-Defined Switching Metrics

Goal: Allow users to define custom policies for runtime adaptation

Custom metric system:

```
class CustomSwitchingPolicy:  
    def should_switch_precision(self, layer_name, current_metrics):  
        latency = current_metrics['latency']  
        accuracy = current_metrics['accuracy']  
  
        if layer_name.startswith('critical_'):  
            return 'float32' if accuracy < 0.99 else 'int16'  
  
        return 'int8' if latency > 50 else 'int16'  
  
    def should_switch_implementation(self, op_name, system_state):  
        if system_state['temperature'] > 70:  
            return 'low_power'  
        return 'high_performance'
```

Register policy:

```
compile_model(model, switching_policy=CustomSwitchingPolicy())
```

Benefits and Challenges

Benefits:

- **Adaptability:** Model adapts to runtime conditions

Challenges:

Benefits and Challenges

Benefits:

- **Adaptability:** Model adapts to runtime conditions
- **Efficiency:** Optimal resource usage

Challenges:

Benefits and Challenges

Benefits:

- **Adaptability:** Model adapts to runtime conditions
- **Efficiency:** Optimal resource usage
- **Flexibility:** User-defined policies

Challenges:

Benefits and Challenges

Benefits:

- **Adaptability:** Model adapts to runtime conditions
- **Efficiency:** Optimal resource usage
- **Flexibility:** User-defined policies
- **Performance:** Balance speed vs. accuracy

Challenges:

Benefits and Challenges

Benefits:

- **Adaptability:** Model adapts to runtime conditions
- **Efficiency:** Optimal resource usage
- **Flexibility:** User-defined policies
- **Performance:** Balance speed vs. accuracy

Challenges:

- Code size increases (multiple implementations)

Benefits and Challenges

Benefits:

- **Adaptability:** Model adapts to runtime conditions
- **Efficiency:** Optimal resource usage
- **Flexibility:** User-defined policies
- **Performance:** Balance speed vs. accuracy

Challenges:

- Code size increases (multiple implementations)
- Runtime overhead for decision logic

Benefits and Challenges

Benefits:

- **Adaptability:** Model adapts to runtime conditions
- **Efficiency:** Optimal resource usage
- **Flexibility:** User-defined policies
- **Performance:** Balance speed vs. accuracy

Challenges:

- Code size increases (multiple implementations)
- Runtime overhead for decision logic
- Complexity in debugging/testing

Benefits and Challenges

Benefits:

- **Adaptability:** Model adapts to runtime conditions
- **Efficiency:** Optimal resource usage
- **Flexibility:** User-defined policies
- **Performance:** Balance speed vs. accuracy

Challenges:

- Code size increases (multiple implementations)
- Runtime overhead for decision logic
- Complexity in debugging/testing
- Policy design requires domain knowledge

Thank You!

Questions?

<https://github.com/ashitabh8/Tiny-NN-in-C.git>